

rcss : SUBGRADIENT AND DUALITY APPROACH FOR DYNAMIC PROGRAMMING

JURI HINZ AND JEREMY YEE*

ABSTRACT. This short paper gives an introduction to the *rcss* package. The R package *rcss* provides users with a tool to approximate the value functions in the Bellman recursion using convex piecewise linear functions formed using operations on tangents. A pathwise method is then used to gauge the quality of the numerical results.

Keywords. Convexity, Dynamic programming, Duality, Subgradient

1. INTRODUCTION

Sequential decision making is often addressed under the framework of *Markov Decision Processes/Dynamic Programming*. However, deriving analytical solutions for even some of the simplest decision processes may be too cumbersome [13, 1, 12]. The use of numerical approximations may be far more practical given the rapid improvements in everyday computational power. The ability to gauge the quality of these approximations is also of significant practical importance. This paper will describe the implementation of fast and accurate algorithms to address these issues for Markov decision processes within a finite time setting, finite action set, convex reward functions and whose Markov processes follow linear dynamics. Under certain conditions, [5] showed that these value function approximations enjoy uniform convergence on compact sets. The package *rcss* represents a *R* implementation of these methods and has already been used to address problems such as pricing financial options [7], natural resource extraction [6], battery management [9], optimal portfolio liquidation [8] and optimal asset allocation under hidden state dynamics [10]. One of the major benefits of implementing these methods in *R* [14] is that the results can be analysed using the vast number of statistical tools available in this language. The *R* package can be found here: <https://github.com/YeeJeremy/rcss> and the manual is listed at <https://github.com/YeeJeremy/RPackageManuals/blob/master/rcss-manual.pdf>.

E-mail address: jeremyyee@outlook.com.au.

2. PROBLEM SETTING

Suppose that state space $\mathbf{X} = \mathbf{P} \times \mathbf{Z}$ is the product of a finite set \mathbf{P} and an open convex set $\mathbf{Z} \subseteq \mathbb{R}^d$. At each decision time $t \in \{0, 1, \dots, T-1\}$, an action $a \in \mathbf{A}$ is chosen by the agent and the dynamic choice of these actions influences the evolution of the Markov process $(X_t)_{t=0}^T := (P_t, Z_t)_{t=0}^T : \Omega \rightarrow \mathbf{P} \times \mathbf{Z}$ where Ω is the set of sample paths. The discrete component $(P_t)_{t=0}^T$ is assumed to be a controlled Markov chain with transition probabilities $(\alpha_{p,p'}^a)_{p,p' \in \mathbf{P}, a \in \mathbf{A}}$, where $\alpha_{p,p'}^a$ is the probability of transitioning from p to p' after applying action a . The second component $(Z_t)_{t=0}^T$ evolves in a linear fashion given by $Z_{t+1} = W_{t+1}Z_t$ where $(W_t)_{t=1}^T$ are matrix-valued random variables referred to as disturbances. The matrix entries in these disturbances are assumed to be integrable. At each time $t = 0, \dots, T-1$ the decision rule π_t is given by a mapping $\pi_t : \mathbf{X} \rightarrow \mathbf{A}$, prescribing at time t an action $\pi_t(p, z) \in \mathbf{A}$ for a given state $(p, z) \in \mathbf{X}$. A sequence $\pi = (\pi_t)_{t=0}^{T-1}$ of decision rules is called a policy. For each policy $\pi = (\pi_t)_{t=0}^{T-1}$, associate it with a so-called policy value $v_0^\pi(p_0, z_0)$ defined as the total expected reward

$$v_0^\pi(p_0, z_0) = \mathbb{E}^{x_0, \pi} \left[\sum_{t=0}^{T-1} r_t(P_t, Z_t, \pi_t(X_t)) + r_T(P_T, Z_T) \right]$$

where $r_T : \mathbf{P} \times \mathbf{Z} \rightarrow \mathbb{R}$ and $r_t : \mathbf{P} \times \mathbf{Z} \times \mathbf{A} \rightarrow \mathbb{R}$ are convex functions in the second argument for $t = 0, \dots, T-1$. These functions represent the scrap and reward in the decision problem, respectively. A policy $\pi^* = (\pi_t^*)_{t=0}^{T-1}$ is called optimal if it maximizes the total expected reward over all policies $\pi \mapsto v_0^\pi(p, z)$. To obtain such policy, one introduces for $t = 0, \dots, T-1$ the so-called *Bellman operator*

$$\mathcal{T}_t v(p, z) = \max_{a \in \mathbf{A}} \left\{ r_t(p, z, a) + \sum_{p' \in \mathbf{P}} \alpha_{p,p'}^a \mathbb{E}^W [v(p', W_{t+1}z)] \right\}, \quad (p, z) \in \mathbf{P} \times \mathbf{Z}$$

acting on all functions v where the expectation is defined. Consider the Bellman recursion, also referred to as backward induction:

$$v_T(p, z) = r_T(p, z), \quad v_t = \mathcal{T}_t v_{t+1} \quad \text{for } t = T-1, \dots, 0.$$

A recursive solution $(v_t^*)_{t=0}^T$ to the Bellman recursion above are called value functions and they determine an optimal policy $\pi^* = (\pi_t^*)_{t=0}^{T-1}$ via

$$\pi_t^*(p, z) = \arg \max_{a \in \mathbf{A}} \left\{ r_t(p, z, a) + \sum_{p' \in \mathbf{P}} \alpha_{p,p'}^a \mathbb{E}^W [v_{t+1}^*(p', W_{t+1}z)] \right\},$$

for $t = T-1, \dots, 0$.

3. NUMERICAL APPROACH

Since the reward and scrap functions are convex in the continuous variable, the value functions are also convex due to the linear state dynamics and so can be approximated by convex piecewise linear functions. For this, introduce the so-called subgradient envelope $\mathcal{S}_{\mathbf{G}^m} f$ of a convex function $f : \mathbf{Z} \rightarrow \mathbb{R}$ on a grid $\mathbf{G}^m \subset \mathbf{Z}$ with m points i.e. $\mathbf{G}^m = \{g^1, \dots, g^m\}$ by

$$\mathcal{S}_{\mathbf{G}^m} f = \bigvee_{g \in \mathbf{G}^m} (\nabla_g f)$$

which is a maximum of the tangents $\nabla_g f$ of f on all grid points $g \in \mathbf{G}^m$. Using the subgradient envelope operator, define the double-modified Bellman operator as

$$\mathcal{T}_t^{m,n} v(p, \cdot) = \mathcal{S}_{\mathbf{G}^m} \max_{a \in \mathbf{A}} \left(r_t(p, \cdot, a) + \sum_{p' \in \mathbf{P}} \alpha_{p,p'}^a \sum_{k=1}^n \nu_{t+1}^{(k)} v(p', W_{t+1}^{(k)} \cdot) \right)$$

where the probability weights $(\nu_{t+1}^{(k)})_{k=1}^n$ corresponds to the distribution sampling $(W_{t+1}^{(k)})_{k=1}^n$ of each disturbance W_{t+1} . The corresponding backward induction

$$\begin{aligned} v_{T-1}^{m,n}(p, z) &= \mathcal{T}_{T-1}^{m,n} \mathcal{S}_{\mathbf{G}^m} r_T(p, z), \\ v_t^{m,n}(p, z) &= \mathcal{T}_t^{m,n} v_{t+1}^{m,n}(p, z), \quad t = T-2, \dots, 0. \end{aligned}$$

for $p \in \mathbf{P}$ and $z \in \mathbf{Z}$ yields the so-called double-modified value functions $(v_t^{m,n})_{t=0}^T$. If the disturbance sampling is constructed using local averages on a partition of the disturbance space or using random Monte Carlo sampling, it can be shown that the double-modified value functions converge uniformly to the true value functions on compact sets if the grid becomes dense in \mathbf{Z} . Now, to gauge the quality of the approximations from the above, we construct two random variables whose expectations bound the true value function i.e.

$$(1) \quad \mathbb{E}(\underline{v}_0(p, z_0)) \leq v_0(p, z_0) \leq \mathbb{E}(\bar{v}_0(p, z_0)), \quad p \in \mathbf{P}, \quad z_0 \in \mathbf{Z}.$$

This process exhibits a helpful self-tuning property. The the closer the value function approximations to optimality, the tighter the bounds in Equation 1 and the lower the standard errors of the bound estimates.

The R package *rcss* represents these convex piecewise linear functions as matrices and offers several options to use nearest neighbour algorithms (from [11]) to reduce the computational cost of the above methods. Most of the computational work is done in C++ via *Rcpp* [3] and is parallelized using *OpenMp* [2]. The following sections will demonstrate some real world applications of this R package.

4. EXAMPLE: BERMUDA PUT

Optimal switching problems naturally arise in the valuation of financial contracts. A simple example is given by the Bermudan Put option. This option gives its owner the right but not an obligation to choose a time to exercise the option in order to receive a payment which depends on the price of the underlying asset at the exercise time. The so-called fair price of the Bermudan option is related to the solution of an optimal stopping problem (see [4]). Here, the asset price process $(\tilde{Z}_t)_{t=0}^T$ at time steps $0, \dots, T$ is modelled as a sampled geometric Brownian motion

$$\tilde{Z}_{t+1} = \epsilon_{t+1} \tilde{Z}_t, \quad t = 0, \dots, T-1, \quad Z_0 \in \mathbb{R}_+,$$

where $(\epsilon_t)_{t=1}^T$ are independent random variables following a log-normal distribution. The fair price of such option with strike price K , interest rate $\rho \geq 0$ and maturity date T , is given by the solution to the optimal stopping problem

$$\sup\{\mathbb{E}(\max(e^{-\rho\tau}(K - \tilde{Z}_\tau), 0)) : \tau \text{ is } \{0, 1, \dots, T\}\text{-valued stopping time}\}.$$

A transformation of the state space is required to be able representing the reward functions in a convenient way for the `rcss` package to process, thus we introduce an augmentation with 1 via

$$Z_t = \begin{bmatrix} 1 \\ \tilde{Z}_t \end{bmatrix}, \quad t = 0, \dots, T.$$

then it becomes possible to represent the evolution as the linear state dynamics

$$Z_{t+1} = W_{t+1} Z_t, \quad t = 0, \dots, T-1$$

with independent and identically distributed matrix-valued random variables $(W_t)_{t=1}^T$ given by

$$W_{t+1} = \begin{bmatrix} 1 & 0 \\ 0 & \epsilon_{t+1} \end{bmatrix}, \quad t = 0, \dots, T-1.$$

This switching system is defined by two positions $\mathbf{P} = \{1, 2\}$ and two actions $\mathbf{A} = \{1, 2\}$. Here, the positions ‘exercised’ and ‘not exercised’ are represented by $p = 1$, $p = 2$ respectively, and the actions ‘don’t exercise’ and ‘exercise’ are denoted by $a = 1$ and $a = 2$ respectively. With this interpretation, the position change is given by deterministic transitions to specified states

$$\alpha_{p,p'}^a = \begin{cases} 1 & \text{if } p' = \alpha(p, a) \\ 0 & \text{else} \end{cases}$$

deterministically determined by the target positions

$$(\alpha(p, a))_{p,a=1}^2 \sim \begin{bmatrix} \alpha(1, 1) & \alpha(1, 2) \\ \alpha(2, 1) & \alpha(2, 2) \end{bmatrix} = \begin{bmatrix} 1 & 1 \\ 2 & 1 \end{bmatrix},$$

while the rewards at time $t = 0, \dots, T$ and are defined as

$$\begin{aligned} r_t(p, (z^{(1)}, z^{(2)}), a) &= e^{-\rho t} \max(K - z^{(2)}, 0)(p - \alpha(p, a)), \\ r_T(p, (z^{(1)}, z^{(2)})) &= e^{-\rho T} \max(K - z^{(2)}, 0)(p - \alpha(p, 2)), \end{aligned}$$

for all $p \in \mathbf{P}$, $a \in \mathbf{A}$, $z \in \mathbb{R}_+$.

4.1. Code Example. As a demonstration, let us consider a Bermuda put option with strike price 40 that expires in 1 year. The put option is exercisable at 51 evenly spaced time points in the year, which includes the start and end of the year. The following code approximates the value functions in the Bellman recursion. On a Linux Ubuntu 16.04 with Intel i5-5300U CPU @2.30GHz and 16GB of RAM, the following code takes around 0.2 cpu second and around 0.05 real world seconds.

LISTING 1. Value function approximation

```

1 library(rcss)
2 rate <- 0.06 ## Interest rate
3 step <- 0.02 ## Time step between decision epochs
4 vol <- 0.2 ## Volatility of stock price process
5 n_dec <- 51 ## Number of decision epochs
6 strike <- 40 ## Strike price
7 control <- matrix(c(c(1, 1), c(2, 1)), nrow = 2, byrow = TRUE) ## Control
8 grid <- as.matrix(cbind(rep(1, 301), seq(30, 60, length = 301))) ## Grid
9 ## Disturbance sampling
10 u <- (rate - 0.5 * vol^2) * step
11 sigma <- vol * sqrt(step)
12 condExpected <- function(a, b){
13   aa <- (log(a) - (u + sigma^2)) / sigma
14   bb <- (log(b) - (u + sigma^2)) / sigma
15   return(exp(u + sigma^2 / 2) * (pnorm(bb) - pnorm(aa)))
16 }
17 weight <- rep(1 / 1000, 1000)
18 disturb <- array(0, dim = c(2, 2, 1000))
19 disturb[1,1,] <- 1
20 part <- qlnorm(seq(0, 1, length = 1000 + 1), u, sigma)
21 for (i in 1:1000) {
22   disturb[2,2,i] <- condExpected(part[i], part[i+1]) / (plnorm(part[i+1],
23     u, sigma) - plnorm(part[i], u, sigma))
24 }
25 ## Subgradient representation of reward
26 in_money <- grid[,2] <= strike
27 reward <- array(0, dim = c(301, 2, 2, 2, n_dec - 1))
28 reward[in_money,1,2,2,] <- strike
29 reward[in_money,2,2,2,] <- -1
30 for (tt in 1:n_dec - 1){
31   reward[,,,,tt] <- exp(-rate * step * (tt - 1)) * reward[,,,,tt]
32 }
33 ## Subgrad representation of scrap
34 scrap <- array(data = 0, dim = c(301, 2, 2))
35 scrap[in_money,1,2] <- strike
36 scrap[in_money,2,2] <- -1

```

```

36 scrap <- exp(-rate * step * (n_dec - 1)) * scrap
37 ## Bellman
38 r_index <- matrix(c(2, 2), ncol = 2)
39 bellman <- FastBellman(grid, reward, scrap, control, disturb, weight,
    r_index)

```

The matrix `grid` represents our choice of grid points where each row represents a point. The 3-dimensional array `disturb` represents our sampling of the disturbances where `disturb[, , i]` gives the i -th sample. Here, we use local averages on a 1000 component partition of the disturbance space. The 5-dimensional array `reward` represents the subgradient approximation with `reward[, , a, p, t]` representing $\mathcal{S}_{G^m} r_t(p, ., a)$. The object `bellman` is a list containing the approximations of the value functions and expected value functions for all positions and decision epochs. Please refer to the package manual for the format of the inputs and outputs. To obtain the value function of the Bermuda put option, simply run the plot command below.

LISTING 2. Option value function

```

40 plot(grid[,2], rowSums(bellman$value[, , 2, 1] * grid), type = "l", xlab =
    "Stock Price", ylab = "Option Value")

```




FIGURE 1. Bermuda put value function.

The following code then computes the lower and upper bound estimates for the value of the option when $\tilde{Z}_0 = 36$. On our machine, the following takes around 10 cpu seconds and around 5 real world seconds to run.

LISTING 3. Lower and upper bounds

```

41 ## Reward function
42 RewardFunc <- function(state, time) {
43   output <- array(data = 0, dim = c(nrow(state), 2, 2))
44   output[,2,2] <- exp(-rate * step * (time - 1)) * pmax(40 - state[,2], 0)
45   return(output)
46 }
47 ## Scrap function
48 ScrapFunc <- function(state) {
49   output <- array(data = 0, dim = c(nrow(state), 2))
50   output[,2] <- exp(-rate * step * (n_dec - 1)) * pmax(40 - state[,2], 0)
51   return(output)
52 }
53 ## Get primal-dual bounds
54 start <- c(1, 36)
55 ## Path disturbances
56 set.seed(12345)
57 n_path <- 500
58 path_disturb <- array(0, dim = c(2, 2, n_path, n_dec - 1))
59 path_disturb[1, 1,,] <- 1
60 rand1 <- rnorm(n_path * (n_dec - 1) / 2)
61 rand1 <- as.vector(rbind(rand1, -rand1))
62 path_disturb[2, 2,,] <- exp((rate - 0.5 * vol^2) * step + vol * sqrt(step) *
63   rand1)
64 path <- PathDisturb(start, path_disturb)
65 policy <- FastPathPolicy(path, grid, control, RewardFunc, bellman$expected)
66 ## Subsim disturbances
67 n_subsim <- 500
68 subsim <- array(0, dim = c(2, 2, n_subsim, n_path, (n_dec - 1)))
69 subsim[1,1,,,] <- 1
70 rand2 <- rnorm(n_subsim * n_path * (n_dec - 1) / 2)
71 rand2 <- as.vector(rbind(rand2, -rand2))
72 subsim[2,2,,,] <- exp((rate - 0.5 * vol^2) * step + vol * sqrt(step) * rand2)
73 subsim_weight <- rep(1 / n_subsim, n_subsim)
74 mart <- FastAddDual(path, subsim, subsim_weight, grid, bellman$value,
75   ScrapFunc)
76 bounds <- AddDualBounds(path, control, RewardFunc, ScrapFunc, mart, policy)

```

The above code takes the exact reward and scrap functions as inputs. The function `FastPathPolicy` computes the candidate optimal policy. The object `bounds` is a list containing the primals $\underline{v}_t^i(p, z_t)$ and duals $\underline{v}_t^i(p, z_t)$ for each sample path i and each position p at each decision time t . Again, please refer to the package manual for the format of the inputs and outputs. If the price of the underlying asset is 36, the 99% confidence interval for the option price is given by the following.

LISTING 4. 99% confidence interval

```

75 > print(GetBounds(bounds, 0.01, 2))
76 [1] 4.475802 4.480533

```

The package '`rcss`' also allows the user to test the prescribed policy from the Bellman recursion on any supplied set of sample paths. The resulting output can then be further studied with time series analysis or other statistical work. In the following code, we will use the previously generated 500 sample paths to backtest our policy and generate histograms.

LISTING 5. Backtesting Policy

```

77 test <- FullTestPolicy(2, path, control, RewardFunc, ScrapFunc, policy)
78 ## Histogram of cumulated rewards
79 hist(test$value, xlab = "Cumulated Rewards", main = "")
80 ## Exercise times
81 ex <- apply(test$position == 1, 1, function(x) min(which(x)))
82 ex[ex == Inf] <- 51
83 ex <- ex - 1
84 hist(ex, xlab = "Exercise Times", main = "")

```




FIGURE 2. Distribution of cumulated rewards and exercise times.

Figure 2 contains the histograms for the cumulated rewards and exercise times. Let us emphasise the usefulness of such scenario generation. Given an approximately optimal policy and backtesting, one can perform statistical analysis on the backtested values to obtain practical insights such as for risk analysis purposes.

5. EXAMPLE: SWING OPTION

Let us now consider the swing option which is a financial contract popular in the energy business. In the simplest form, it gives the owner the right to obtain a certain commodity (such as gas or electricity) at a pre-specified price and volume at a number of exercise times which can be freely chosen by the contract owner. Let us consider a specific case of such contract, referred to as a unit-time refraction period swing option. In this contract, there is a limit to exercise only one right at any time. Given the discounted commodity price $(S_t)_{t=0}^T$, the so-called fair price of a swing option with N rights is given by the supremum

$$\sup_{0 \leq \tau_1 < \dots < \tau_N \leq T} \mathbb{E} \left[\sum_{n=1}^N (S_{\tau_n} - K e^{-\rho \tau_n})^+ \right]$$

over all stopping times τ_1, \dots, τ_N with values in $\{0, \dots, T\}$. In order to represent this control problem as a switching system, we use the position set $\mathbf{P} = \{1, \dots, N+1\}$ to describe the number of exercise rights remaining. That is $p \in \mathbf{P}$ stands for the situation when there are $p - 1$ rights remaining to be exercised. The action set $\mathbf{A} = \{1, 2\}$ represents the choice between exercising ($a = 1$) or not exercising ($a = 2$). The control matrices $(\alpha_{p,p'}^a)$ are given for exercise action $a = 1$

$$\alpha_{p,p'}^1 = \begin{cases} 1 & \text{if } p' = 1 \vee (p - 1) \\ 0 & \text{else,} \end{cases}$$

and for not-exercise action $a = 2$ as

$$\alpha_{p,p'}^2 = \begin{cases} 1 & \text{if } p' = p \\ 0 & \text{else} \end{cases}$$

for all $p, p' \in \mathbf{P}$. In the case of the swing option, the transition between p and p' occurs deterministically, since once the controller decides to exercise the right, the number of rights remaining is diminished by one. The deterministic control of the discrete component is easier to describe in term of the matrix $(\alpha(p, a))_{p \in \mathbf{P}, a \in \mathbf{A}}$ where $p' = \alpha(p, a) \in \mathbf{P}$ stands for the discrete component which is reached from $p \in \mathbf{P}$ by the action $a \in \mathbf{A}$. For the case of the swing option this matrix is

$$(\alpha(p, a))_{p \in \mathbf{P}, a \in \mathbf{A}} = \begin{bmatrix} 1 & 1 \\ 1 & 2 \\ 2 & 3 \\ \dots & \dots \\ N & N+1 \end{bmatrix}.$$

Having modelled the discounted commodity price process as an exponential mean-reverting process with a reversion parameter $\kappa \in [0, 1[$, long run mean

$\mu > 0$ and volatility $\sigma > 0$, we obtain the logarithm of the discounted price process as

$$\tilde{Z}_{t+1} = (1 - \kappa)(\tilde{Z}_t - \mu) + \mu + \sigma\epsilon_{t+1}, \quad \tilde{Z}_0 = \ln(S_0).$$

A further transformation of the state space is required before linear state dynamics can be achieved. If we introduce an augmentation with 1 via

$$Z_t = \begin{bmatrix} 1 \\ \tilde{Z}_t \end{bmatrix}, \quad t = 0, \dots, T.$$

then it becomes possible to represent the evolution as the linear state dynamics

$$Z_{t+1} = W_{t+1}Z_t, \quad t = 0, \dots, T - 1$$

with independent and identically distributed matrix-valued random variables $(W_t)_{t=1}^T$ given by

$$W_{t+1} = \begin{bmatrix} 1 & 0 \\ \kappa\mu + \sigma\epsilon_{t+1} & (1 - \kappa) \end{bmatrix}, \quad t = 0, \dots, T - 1.$$

The reward and scrap values are given by

$$(2) \quad r_t(p, (z^{(1)}, z^{(2)}), a) = (e^{z^{(2)}} - Ke^{-\rho t})^+(p - \alpha(p, a))$$

for $t = 0, \dots, T - 1$ and

$$(3) \quad r_T(p, (z^{(1)}, z^{(2)})) = (e^{z^{(2)}} - Ke^{-\rho T})^+(p - \alpha(p, 1))$$

respectively for all $p \in \mathbf{P}$ and $a \in \mathbf{A}$.

5.1. Code Example. In this example, consider a swing option with 5 rights exercisable on 101 time points. As before, we begin by performing the value function approximation. On our machine, the following code takes around 0.4 cpu seconds or around 0.15 real world seconds to run.

LISTING 6. Value function approximation

```

85 library(rcss)
86 ## Parameters
87 rho <- 0
88 kappa <- 0.9
89 mu <- 0
90 sigma <- 0.5
91 K <- 0
92 n_dec <- 101 ## number of time epochs
93 N <- 5 ## number of rights
94 n_pos <- N + 1 ## number of positions
95 grid <- cbind(rep(1, 101), seq(-2, 2, length = 101)) ## Grid
96 ## Control matrix
97 control <- cbind(c(1, 1:N), 1:(N + 1))
98 ## Reward subgradient representation
99 reward <- array(0, dim = c(101, 2, 2, nrow(control), n_dec - 1))
100 slope <- exp(grid[, 2])

```

```

101 for (tt in 1:(n_dec - 1)) {
102   discount <- exp(-rho * (tt - 1))
103   for (pp in 2:n_pos) {
104     intercept <- (exp(grid[,2]) - K * discount) - slope * grid[, 2]
105     reward[, 1, 1, pp, tt] <- intercept
106     reward[, 2, 1, pp, tt] <- slope
107   }
108 }
109 ## Scrap subgradient representation
110 scrap <- array(0, dim = c(101, 2, nrow(control)))
111 discount <- exp(-rho * (n_dec - 1))
112 for (pp in 2:n_pos) {
113   intercept <- (exp(grid[,2]) - K * discount) - slope * grid[, 2]
114   scrap[, 1, pp] <- intercept
115   scrap[, 2, pp] <- slope
116 }
117 ## Disturbance sampling
118 weight <- rep(1/1000, 1000)
119 disturb <- array(0, dim = c(2, 2, 1000))
120 disturb[1, 1,] <- 1
121 disturb[2, 2,] <- 1 - kappa
122 CondExpected <- function(a, b){
123   return(1/sqrt(2 * pi) * (exp(-a^2/2) - exp(-b^2/2)))
124 }
125 part <- qnorm(seq(0, 1, length = 1000 + 1))
126 for (i in 1:1000) {
127   disturb[2,1,i] <- kappa * mu + sigma * (CondExpected(part[i], part[i+1])
128     / (pnorm(part[i+1]) - pnorm(part[i])))
129 }
129 ## Bellman recursion
130 r_index <- matrix(c(2, 1), ncol = 2)
131 bellman <- FastBellman(grid, reward, scrap, control, disturb, weight,
132   r_index)

```

After obtaining these function approximations, the following code computes the 99% confidence intervals for the value of a swing option with 5 remaining rights. The code below takes approximately 20 cpu seconds or 10 real world seconds to run.

LISTING 7. Lower and upper bounds

```

132 ## Exact reward function
133 RewardFunc <- function(state, time) {
134   output <- array(0, dim = c(nrow(state), 2, nrow(control)))
135   discount <- exp(-rho * (time - 1))
136   for (i in 2:nrow(control)) {
137     output[, 1, i] <- pmax(exp(state[, 2]) - K * discount, 0)
138   }
139   return(output)
140 }
141 ## Exact scrap function
142 ScrapFunc <- function(state) {
143   output <- array(0, dim = c(nrow(state), nrow(control)))
144   discount <- exp(-rho * (n_dec - 1))

```

```

145     for (i in 2:nrow(control)) {
146         output[, i] <- pmax(exp(state[, 2]) - K * discount, 0)
147     }
148     return(output)
149 }
150 ## Generate paths
151 set.seed(12345)
152 n_path <- 500
153 path_disturb <- array(0, dim = c(2, 2, n_path, n_dec - 1))
154 path_disturb[1, 1,,] <- 1
155 path_disturb[2, 2,,] <- 1 - kappa
156 rand1 <- rnorm(n_path * (n_dec - 1) / 2)
157 rand1 <- as.vector(rbind(rand1, -rand1))
158 path_disturb[2, 1,,] <- kappa * mu + sigma * rand1
159 start <- c(1, 0)
160 path <- PathDisturb(start, path_disturb)
161 policy <- FastPathPolicy(path, grid, control, RewardFunc, bellman$expected)
162 ## Set subsimulation disturbances
163 n_subsim <- 500
164 subsim <- array(0, dim = c(2, 2, n_subsim, n_path, n_dec - 1))
165 subsim[1, 1,,,] <- 1
166 subsim[2, 2,,,] <- 1 - kappa
167 rand2 <- rnorm(n_subsim * n_path * (n_dec - 1) / 2)
168 rand2 <- as.vector(rbind(rand2, -rand2))
169 subsim[2, 1,,,] <- kappa * mu + sigma * rand2
170 subsim_weight <- rep(1 / n_subsim, n_subsim)
171 ## Primal-dual
172 mart <- FastAddDual(path, subsim, subsim_weight, grid, bellman$value,
173                     ScrapFunc)
173 bounds <- AddDualBounds(path, control, RewardFunc, ScrapFunc, mart, policy)

```

LISTING 8. 99% confidence interval

```

174 > print(GetBounds(bounds, 0.01, 6))
175 [1] 13.42159 13.44162

```

6. CONCLUSION

This paper gives a demonstration of the R package *rcss* in solving optimal switching problems. The problem setting discussed in this paper is broad and can be used to model a wide range of problems. Using nearest neighbour algorithms, the package *rcss* is able to solve some real world problems in an accurate and quick manner.

REFERENCES

- [1] N. Bauerle and U. Rieder, *Markov decision processes with applications to finance*, Springer, Heidelberg, 2011.
- [2] L. Dagum and R. Menon, *Openmp: an industry standard api for shared-memory programming*, IEEE Computational Science & Engineering **5** (1998), no. 1, 46–55.

- [3] D. Eddelbuettel and R. Francois, *Rcpp: Seamless R and C++ integration*, Journal of Statistical Software **40** (2011), no. 8, 1–18.
- [4] P. Glasserman, *Monte Carlo methods in financial engineering*, Springer, 2003.
- [5] J. Hinz, *Optimal stochastic switching under convexity assumptions*, SIAM Journal on Control and Optimization **52** (2014), no. 1, 164–188.
- [6] J. Hinz, T. Tarnopolskaya, and J. Yee, *Commodity resource valuation and extraction: A pathwise programming approach*, Preprint (Preprint).
- [7] J. Hinz and N. Yap, *Algorithms for optimal control of stochastic switching systems*, Theory of Probability and its Applications **60** (2015), no. 4, 770–800.
- [8] J. Hinz and J. Yee, *An algorithmic approach to optimal asset liquidation problems*, Asia-Pacific Financial Markets **24** (2017), no. 2, 109–129.
- [9] ———, *Optimal forward trading and battery control under renewable electricity generation*, Journal of Banking & Finance **InPress** (2017).
- [10] ———, *Stochastic switching for partially observable dynamics and optimal asset allocation*, International Journal of Control **90** (2017), no. 3, 553–565.
- [11] M. Muja and D. Lowe, *Flann - fast library for approximate nearest neighbors*, 2016, Version 1.9.1.
- [12] H. Pham, *Continuous-time stochastic control and optimization with financial applications*, vol. 61, Springer Science & Business Media, 2009.
- [13] W. Powell, *Approximate dynamic programming: Solving the curses of dimensionality*, Wiley, Hoboken, New Jersey, 2007.
- [14] R Core Team, *R: A language and environment for statistical computing*, R Foundation for Statistical Computing, 2013, ISBN 3-900051-07-0.